

# CNN Assignment : Apply 3 different CNN's on the MNIST dataset

```
In [1]: # Credits: https://github.com/keras-team/keras/blob/master/examples/mnist\_cnn.py
#Refer this link for making better CNN networks
#https://towardsdatascience.com/a-guide-to-an-efficient-way-to-build-neural-network-architecturespart-ii-hyper-parameter-42efca01e
import warnings
warnings.filterwarnings("ignore")
#from __future__ import print_function
exec('from __future__ import absolute_import, division, print_function')
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
batch_size = 128
num_classes = 10
epochs = 12
# Preparing training and testing data
# input image dimensions
img_rows, img_cols = 28, 28
# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()
#print(x_train.shape)
if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
# convert class vectors to binary class matrices
```

```
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

Using TensorFlow backend.

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>

11493376/11490434 [=====] - 3s 0us/step

x\_train shape: (60000, 28, 28, 1)

60000 train samples

10000 test samples

```
In [2]: %matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334 # this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

## Models with Conv , Max Pool and Dense Layer

### Model 1 : 2 conv + 2 maxpoll+ 3 dense layers

```
In [3]: import warnings
warnings.filterwarnings("ignore")
# In this (First Model) lets follow the general structure of the Lenet we will make a simple model
# Network Architecture
# input -> conv -> polling -> conv -> polling -> FC -> FC -> output
# 8 16 120 84 10
model = Sequential()
model.add(Conv2D(8, kernel_size=(3, 3), activation='relu', padding='same', input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Conv2D(16, (5, 5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Flatten())
model.add(Dense(120, activation='relu'))
```

```

model.add(Dense(84, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.adam(),
              metrics=['accuracy'])
# this will train the model and validate the model in this fit function
model.summary()

```

WARNING:tensorflow:From C:\Anaconda\lib\site-packages\tensorflow\python\ops\resource\_variable\_ops.py:435: colocate\_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 28, 28, 8)	80
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 8)	0
conv2d_2 (Conv2D)	(None, 10, 10, 16)	3216
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 16)	0
flatten_1 (Flatten)	(None, 400)	0
dense_1 (Dense)	(None, 120)	48120
dense_2 (Dense)	(None, 84)	10164
dense_3 (Dense)	(None, 10)	850
Total params: 62,430		
Trainable params: 62,430		
Non-trainable params: 0		

```

In [4]: import warnings
warnings.filterwarnings("ignore")
history=model.fit(x_train, y_train,
                 batch_size=batch_size,
                 epochs=epochs,
                 verbose=1,
                 validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)

```

```
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

WARNING:tensorflow:From C:\Anaconda\lib\site-packages\tensorflow\python\ops\math\_ops.py:3066: to\_int32 (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 17s 287us/step - loss: 0.2721 - accuracy: 0.9219 - val\_loss: 0.0698 - val\_accuracy: 0.9765

Epoch 2/12

60000/60000 [=====] - 18s 298us/step - loss: 0.0740 - accuracy: 0.9771 - val\_loss: 0.0515 - val\_accuracy: 0.9838

Epoch 3/12

60000/60000 [=====] - 18s 298us/step - loss: 0.0528 - accuracy: 0.9838 - val\_loss: 0.0404 - val\_accuracy: 0.9859

Epoch 4/12

60000/60000 [=====] - 18s 306us/step - loss: 0.0425 - accuracy: 0.9868 - val\_loss: 0.0338 - val\_accuracy: 0.9889

Epoch 5/12

60000/60000 [=====] - 20s 332us/step - loss: 0.0353 - accuracy: 0.9890 - val\_loss: 0.0345 - val\_accuracy: 0.9885

Epoch 6/12

60000/60000 [=====] - 18s 300us/step - loss: 0.0294 - accuracy: 0.9903 - val\_loss: 0.0421 - val\_accuracy: 0.9864

Epoch 7/12

60000/60000 [=====] - 18s 293us/step - loss: 0.0252 - accuracy: 0.9920 - val\_loss: 0.0321 - val\_accuracy: 0.9888

Epoch 8/12

60000/60000 [=====] - 18s 296us/step - loss: 0.0205 - accuracy: 0.9931 - val\_loss: 0.0373 - val\_accuracy: 0.9887

Epoch 9/12

60000/60000 [=====] - 17s 289us/step - loss: 0.0184 - accuracy: 0.9937 - val\_loss: 0.0352 - val\_accuracy: 0.9888

Epoch 10/12

60000/60000 [=====] - 18s 299us/step - loss: 0.0172 - accuracy: 0.9942 - val\_loss: 0.0415 - val\_accuracy: 0.9859

Epoch 11/12

60000/60000 [=====] - 18s 298us/step - loss: 0.0136 - accuracy: 0.9957 - val\_loss: 0.0342 - val\_accuracy: 0.9892

Epoch 12/12

60000/60000 [=====] - 18s 298us/step - loss: 0.0132 - accuracy: 0.9956 - val\_loss: 0.0464 - val\_accuracy: 0.9871

Test loss: 0.046447735224399364

Test accuracy: 0.9871000051498413

```

In [5]: score = model.evaluate(x_train, y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n*****\n')
#test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
# plot
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch');
ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,12+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```


Train score: 0.01593446881301449

Train accuracy: 99.4533360004425

\*\*\*\*\*

Test score: 0.046447735224399364

Test accuracy: 98.71000051498413



## Model 2 : 3 conv + 3 maxpoll+ 2 dense layers

```
In [6]: import warnings
warnings.filterwarnings("ignore")
# go basic model to deep layer model
# Network Architecture
# input -> conv -> pooling -> conv -> pooling -> conv -> pooling -> FC -> output
# 8 32 128 64
model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Conv2D(64, (3,3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Flatten())
model.add(Dense(64, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
optimizer=keras.optimizers.adam(),
metrics=['accuracy'])
# this will train the model and validate the model in this fit function
model.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_3 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_4 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_4 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_5 (Conv2D)	(None, 3, 3, 128)	73856
max_pooling2d_5 (MaxPooling2D)	(None, 1, 1, 128)	0
flatten_2 (Flatten)	(None, 128)	0
dense_4 (Dense)	(None, 64)	8256
dense_5 (Dense)	(None, 10)	650
=====		
Total params: 101,578		
Trainable params: 101,578		

Non-trainable params: 0

```
In [7]: import warnings
warnings.filterwarnings("ignore")
history=model.fit(x_train, y_train,
                  batch_size=batch_size,
                  epochs=epochs,
                  verbose=1,
                  validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 48s 796us/step - loss: 0.3081 - accuracy: 0.9097 - val\_loss: 0.0952 - val\_accuracy: 0.9729

Epoch 2/12

60000/60000 [=====] - 46s 766us/step - loss: 0.0917 - accuracy: 0.9721 - val\_loss: 0.0805 - val\_accuracy: 0.9736

Epoch 3/12

60000/60000 [=====] - 45s 758us/step - loss: 0.0675 - accuracy: 0.9797 - val\_loss: 0.0624 - val\_accuracy: 0.9813

Epoch 4/12

60000/60000 [=====] - 45s 750us/step - loss: 0.0539 - accuracy: 0.9836 - val\_loss: 0.0514 - val\_accuracy: 0.9841

Epoch 5/12

60000/60000 [=====] - 45s 757us/step - loss: 0.0433 - accuracy: 0.9865 - val\_loss: 0.0573 - val\_accuracy: 0.9827

Epoch 6/12

60000/60000 [=====] - 45s 757us/step - loss: 0.0359 - accuracy: 0.9889 - val\_loss: 0.0682 - val\_accuracy: 0.9794

Epoch 7/12

60000/60000 [=====] - 45s 755us/step - loss: 0.0318 - accuracy: 0.9900 - val\_loss: 0.0545 - val\_accuracy: 0.9834

Epoch 8/12

60000/60000 [=====] - 46s 764us/step - loss: 0.0269 - accuracy: 0.9910 - val\_loss: 0.0499 - val\_accuracy: 0.9866

Epoch 9/12

60000/60000 [=====] - 46s 761us/step - loss: 0.0228 - accuracy: 0.9928 - val\_loss: 0.0533 - val\_accuracy: 0.9857

Epoch 10/12

60000/60000 [=====] - 46s 759us/step - loss: 0.0197 - accuracy: 0.9937 - val\_loss: 0.0485 - val\_accuracy: 0.9868

Epoch 11/12

60000/60000 [=====] - 49s 823us/step - loss: 0.0169 - accuracy: 0.9944 - val\_loss: 0.0503 - val\_accuracy:

```

0.9874
Epoch 12/12
60000/60000 [=====] - 48s 793us/step - loss: 0.0150 - accuracy: 0.9950 - val_loss: 0.0501 - val_accuracy:
0.9851
Test loss: 0.05009871911372902
Test accuracy: 0.9850999712944031

```

```

In [8]: score = model.evaluate(x_train, y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n*****\n')
#test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
# plot
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch');
ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,12+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Train score: 0.011942215900942877
Train accuracy: 99.6150016784668

```

```

*****


```

```

Test score: 0.05009871911372902
Test accuracy: 98.50999712944031

```





Finally we train a model with the trend Conv-Conv-Pool-Conv-Conv-Pool

## Model 3 : 4 conv+ 2 maxpoll + 2 dense Layer

```
In [9]: # go basic model to deep layer model
# Network Architecture
# input -> conv -> conv -> polling -> conv -> conv -> polling -> FC -> output
# 16 16 32 32 512
model = Sequential()
model.add(Conv2D(16, kernel_size=(3, 3), activation='relu', padding='same', input_shape=input_shape))
model.add(Conv2D(16, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2)) # for the location invariants
model.add(Conv2D(32, (3, 3), activation='relu'))
model.add(Conv2D(32, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2)) # for the location invariants
model.add(Flatten())
model.add(Dense(512, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
optimizer=keras.optimizers.adam(),
metrics=['accuracy'])
# this will train the model and validate the model in this fit function
model.summary()
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 28, 28, 16)	160
conv2d_7 (Conv2D)	(None, 28, 28, 16)	2320
max_pooling2d_6 (MaxPooling2D)	(None, 14, 14, 16)	0
conv2d_8 (Conv2D)	(None, 12, 12, 32)	4640
conv2d_9 (Conv2D)	(None, 10, 10, 32)	9248
max_pooling2d_7 (MaxPooling2D)	(None, 5, 5, 32)	0
flatten_3 (Flatten)	(None, 800)	0
dense_6 (Dense)	(None, 512)	410112
dense_7 (Dense)	(None, 10)	5130
Total params: 431,610		
Trainable params: 431,610		
Non-trainable params: 0		

```
In [10]: import warnings
warnings.filterwarnings("ignore")
history=model.fit(x_train, y_train,
                  batch_size=batch_size,
                  epochs=epochs,
                  verbose=1,
                  validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 83s 1ms/step - loss: 0.1849 - accuracy: 0.9430 - val\_loss: 0.0541 - val\_accuracy: 0.9824

Epoch 2/12

60000/60000 [=====] - 85s 1ms/step - loss: 0.0485 - accuracy: 0.9848 - val\_loss: 0.0321 - val\_accuracy: 0.9892

Epoch 3/12

60000/60000 [=====] - 83s 1ms/step - loss: 0.0321 - accuracy: 0.9897 - val\_loss: 0.0336 - val\_accuracy: 0.9899

```

Epoch 4/12
60000/60000 [=====] - 84s 1ms/step - loss: 0.0236 - accuracy: 0.9922 - val_loss: 0.0330 - val_accuracy: 0.9887
Epoch 5/12
60000/60000 [=====] - 81s 1ms/step - loss: 0.0188 - accuracy: 0.9942 - val_loss: 0.0259 - val_accuracy: 0.9913
Epoch 6/12
60000/60000 [=====] - 80s 1ms/step - loss: 0.0151 - accuracy: 0.9948 - val_loss: 0.0294 - val_accuracy: 0.9916
Epoch 7/12
60000/60000 [=====] - 84s 1ms/step - loss: 0.0123 - accuracy: 0.9959 - val_loss: 0.0371 - val_accuracy: 0.9890
Epoch 8/12
60000/60000 [=====] - 81s 1ms/step - loss: 0.0106 - accuracy: 0.9963 - val_loss: 0.0303 - val_accuracy: 0.9907
Epoch 9/12
60000/60000 [=====] - 79s 1ms/step - loss: 0.0084 - accuracy: 0.9974 - val_loss: 0.0289 - val_accuracy: 0.9927
Epoch 10/12
60000/60000 [=====] - 90s 2ms/step - loss: 0.0099 - accuracy: 0.9968 - val_loss: 0.0282 - val_accuracy: 0.9929
Epoch 11/12
60000/60000 [=====] - 89s 1ms/step - loss: 0.0089 - accuracy: 0.9970 - val_loss: 0.0301 - val_accuracy: 0.9917
Epoch 12/12
60000/60000 [=====] - 86s 1ms/step - loss: 0.0052 - accuracy: 0.9984 - val_loss: 0.0326 - val_accuracy: 0.9930
Test loss: 0.03256005091774296
Test accuracy: 0.9929999709129333

```

```

In [11]: score = model.evaluate(x_train, y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n*****\n')
#test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
# plot
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch');
ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,12+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```


Train score: 0.003995988390061666

Train accuracy: 99.85666871070862

\*\*\*\*\*

Test score: 0.03256005091774296

Test accuracy: 99.29999709129333



## Models included Dropout

### Model 1 : 2 conv + 2 maxpoll+ 3 dense layer +Dropout (0.5)

```
In [12]: #Same models with Dropouts
import warnings
warnings.filterwarnings("ignore")
# In this (First Model) Lets follow the general structure of the Lenet we will make a simple model
# Network Architecture
# input -> conv -> polling -> conv -> polling ->droupout-> FC -> FC -> output
# 8 16 120 84 10
model = Sequential()
model.add(Conv2D(8, kernel_size=(3, 3),activation='relu',padding='same',input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2),strides=2))# for the location invariants
model.add(Conv2D(16, (5, 5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2),strides=2))# for the location invariants
model.add(Dropout(0.5))
```

```

model.add(Flatten())
model.add(Dense(120, activation='relu'))
model.add(Dense(84, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
optimizer=keras.optimizers.adam(),
metrics=['accuracy'])
# this will train the model and validate the model in this fit function
model.summary()

```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
conv2d_10 (Conv2D)	(None, 28, 28, 8)	80
max_pooling2d_8 (MaxPooling2D)	(None, 14, 14, 8)	0
conv2d_11 (Conv2D)	(None, 10, 10, 16)	3216
max_pooling2d_9 (MaxPooling2D)	(None, 5, 5, 16)	0
dropout_1 (Dropout)	(None, 5, 5, 16)	0
flatten_4 (Flatten)	(None, 400)	0
dense_8 (Dense)	(None, 120)	48120
dense_9 (Dense)	(None, 84)	10164
dense_10 (Dense)	(None, 10)	850
Total params: 62,430		
Trainable params: 62,430		
Non-trainable params: 0		

```

In [13]: history=model.fit(x_train, y_train,
                        batch_size=batch_size,
                        epochs=epochs,
                        verbose=1,
                        validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/12
60000/60000 [=====] - 17s 276us/step - loss: 0.3910 - accuracy: 0.8756 - val_loss: 0.0814 - val_accuracy: 0.9738
Epoch 2/12
60000/60000 [=====] - 18s 305us/step - loss: 0.1317 - accuracy: 0.9588 - val_loss: 0.0547 - val_accuracy: 0.9819
Epoch 3/12
60000/60000 [=====] - 18s 303us/step - loss: 0.1035 - accuracy: 0.9679 - val_loss: 0.0398 - val_accuracy: 0.9868
Epoch 4/12
60000/60000 [=====] - 18s 296us/step - loss: 0.0866 - accuracy: 0.9733 - val_loss: 0.0354 - val_accuracy: 0.9877
Epoch 5/12
60000/60000 [=====] - 18s 296us/step - loss: 0.0780 - accuracy: 0.9757 - val_loss: 0.0317 - val_accuracy: 0.9898
Epoch 6/12
60000/60000 [=====] - 18s 295us/step - loss: 0.0688 - accuracy: 0.9787 - val_loss: 0.0368 - val_accuracy: 0.9881
Epoch 7/12
60000/60000 [=====] - 18s 299us/step - loss: 0.0631 - accuracy: 0.9796 - val_loss: 0.0324 - val_accuracy: 0.9896
Epoch 8/12
60000/60000 [=====] - 18s 298us/step - loss: 0.0587 - accuracy: 0.9812 - val_loss: 0.0268 - val_accuracy: 0.9916
Epoch 9/12
60000/60000 [=====] - 18s 298us/step - loss: 0.0553 - accuracy: 0.9825 - val_loss: 0.0283 - val_accuracy: 0.9910
Epoch 10/12
60000/60000 [=====] - 18s 300us/step - loss: 0.0523 - accuracy: 0.9834 - val_loss: 0.0272 - val_accuracy: 0.9910
Epoch 11/12
60000/60000 [=====] - 18s 298us/step - loss: 0.0515 - accuracy: 0.9837 - val_loss: 0.0269 - val_accuracy: 0.9906
Epoch 12/12
60000/60000 [=====] - 18s 300us/step - loss: 0.0474 - accuracy: 0.9846 - val_loss: 0.0236 - val_accuracy: 0.9912
Test loss: 0.02355184760145494
Test accuracy: 0.9911999702453613

```

```

In [14]: score = model.evaluate(x_train, y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n*****\n')
#test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)

```

```
# plot
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch');
ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,12+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```


Train score: 0.017866180427525736

Train accuracy: 99.43666458129883

\*\*\*\*\*

Test score: 0.02355184760145494

Test accuracy: 99.11999702453613



## Model 2 : 3 conv + 3 maxpoll+ 2 dense layers + Dropout (0.9)

```
In [15]: import warnings
warnings.filterwarnings("ignore")
# go basic model to deep layer model
# Network Architecture
# input -> conv -> polling -> conv -> polling -> conv -> polling -> dropout-> FC -> output
# 8 32 128 64
model = Sequential()
```

```

model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Conv2D(64, (3,3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Dropout(0.9))
model.add(Flatten())
model.add(Dense(64, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
optimizer=keras.optimizers.adam(),
metrics=['accuracy'])
# this will train the model and validate the model in this fit function
model.summary()

```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
=====		
conv2d_12 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_10 (MaxPooling)	(None, 13, 13, 32)	0
conv2d_13 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_11 (MaxPooling)	(None, 5, 5, 64)	0
conv2d_14 (Conv2D)	(None, 3, 3, 128)	73856
max_pooling2d_12 (MaxPooling)	(None, 1, 1, 128)	0
dropout_2 (Dropout)	(None, 1, 1, 128)	0
flatten_5 (Flatten)	(None, 128)	0
dense_11 (Dense)	(None, 64)	8256
dense_12 (Dense)	(None, 10)	650
=====		
Total params: 101,578		
Trainable params: 101,578		
Non-trainable params: 0		

In [16]: `history=model.fit(x_train, y_train,`



```

batch_size=batch_size,
epochs=epochs,
verbose=1,
validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 44s 741us/step - loss: 1.3356 - accuracy: 0.5181 - val\_loss: 0.2980 - val\_accuracy: 0.9365

Epoch 2/12

60000/60000 [=====] - 45s 747us/step - loss: 0.8161 - accuracy: 0.7112 - val\_loss: 0.1750 - val\_accuracy: 0.9537

Epoch 3/12

60000/60000 [=====] - 45s 752us/step - loss: 0.6775 - accuracy: 0.7623 - val\_loss: 0.1483 - val\_accuracy: 0.9616

Epoch 4/12

60000/60000 [=====] - 45s 751us/step - loss: 0.6011 - accuracy: 0.7896 - val\_loss: 0.1250 - val\_accuracy: 0.9640

Epoch 5/12

60000/60000 [=====] - 45s 748us/step - loss: 0.5565 - accuracy: 0.8058 - val\_loss: 0.1123 - val\_accuracy: 0.9697

Epoch 6/12

60000/60000 [=====] - 46s 761us/step - loss: 0.5184 - accuracy: 0.8206 - val\_loss: 0.1050 - val\_accuracy: 0.9705

Epoch 7/12

60000/60000 [=====] - 46s 760us/step - loss: 0.4900 - accuracy: 0.8301 - val\_loss: 0.1024 - val\_accuracy: 0.9713

Epoch 8/12

60000/60000 [=====] - 45s 753us/step - loss: 0.4653 - accuracy: 0.8412 - val\_loss: 0.1005 - val\_accuracy: 0.9722

Epoch 9/12

60000/60000 [=====] - 46s 762us/step - loss: 0.4491 - accuracy: 0.8457 - val\_loss: 0.1117 - val\_accuracy: 0.9685

Epoch 10/12

60000/60000 [=====] - 45s 755us/step - loss: 0.4229 - accuracy: 0.8566 - val\_loss: 0.1161 - val\_accuracy: 0.9675

Epoch 11/12

60000/60000 [=====] - 45s 754us/step - loss: 0.4140 - accuracy: 0.8597 - val\_loss: 0.0912 - val\_accuracy: 0.9745s - los

Epoch 12/12

60000/60000 [=====] - 45s 754us/step - loss: 0.3991 - accuracy: 0.8645 - val\_loss: 0.0948 - val\_accuracy: 0.9748

Test loss: 0.09479323272332549

Test accuracy: 0.9747999906539917

```
In [17]: keras.layers.BatchNormalization(axis=-1, momentum=0.99, epsilon=0.001, center=True, scale=True, beta_initializer='zeros', gamma_in

Out[17]: <keras.layers.normalization.BatchNormalization at 0x1d744fe5a90>
```

## Model 3 : 4 conv + 2 maxpoll+ 2 dense layers + Dropout (0.3)

```
In [18]: # go basic model to deep layer model
# Network Architecture
# input -> conv -> conv -> polling -> conv -> conv -> polling -> dropout-> FC -> output
# 16 16 32 32 512
model = Sequential()
model.add(Conv2D(16, kernel_size=(3, 3),activation='relu',padding='same',input_shape=input_shape))
model.add(Conv2D(16,(3, 3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2),strides=2))# for the location invariants
model.add(Conv2D(32, (3,3), activation='relu'))
model.add(Conv2D(32, (3,3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2),strides=2))# for the location invariants
model.add(Dropout(0.3))
model.add(Flatten())
model.add(Dense(512, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
optimizer=keras.optimizers.adam(),
metrics=['accuracy'])
# this will train the model and validate the model in this fit function
model.summary()
```

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
=====		
conv2d_15 (Conv2D)	(None, 28, 28, 16)	160
conv2d_16 (Conv2D)	(None, 28, 28, 16)	2320
max_pooling2d_13 (MaxPooling)	(None, 14, 14, 16)	0
conv2d_17 (Conv2D)	(None, 12, 12, 32)	4640
conv2d_18 (Conv2D)	(None, 10, 10, 32)	9248
max_pooling2d_14 (MaxPooling)	(None, 5, 5, 32)	0

dropout_3 (Dropout)	(None, 5, 5, 32)	0
flatten_6 (Flatten)	(None, 800)	0
dense_13 (Dense)	(None, 512)	410112
dense_14 (Dense)	(None, 10)	5130
=====		
Total params: 431,610		
Trainable params: 431,610		
Non-trainable params: 0		

```
In [19]: history=model.fit(x_train, y_train,
batch_size=batch_size,
epochs=epochs,
verbose=1,
validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 81s 1ms/step - loss: 0.2103 - accuracy: 0.9357 - val\_loss: 0.0499 - val\_accuracy: 0.9849

Epoch 2/12

60000/60000 [=====] - 80s 1ms/step - loss: 0.0618 - accuracy: 0.9807 - val\_loss: 0.0332 - val\_accuracy: 0.9881

Epoch 3/12

60000/60000 [=====] - 80s 1ms/step - loss: 0.0422 - accuracy: 0.9871 - val\_loss: 0.0282 - val\_accuracy: 0.9905

Epoch 4/12

60000/60000 [=====] - 80s 1ms/step - loss: 0.0350 - accuracy: 0.9891 - val\_loss: 0.0250 - val\_accuracy: 0.9914

Epoch 5/12

60000/60000 [=====] - 80s 1ms/step - loss: 0.0289 - accuracy: 0.9907 - val\_loss: 0.0217 - val\_accuracy: 0.9921

Epoch 6/12

60000/60000 [=====] - 80s 1ms/step - loss: 0.0241 - accuracy: 0.9919 - val\_loss: 0.0277 - val\_accuracy: 0.9923

Epoch 7/12

60000/60000 [=====] - 80s 1ms/step - loss: 0.0217 - accuracy: 0.9929 - val\_loss: 0.0242 - val\_accuracy: 0.9928

Epoch 8/12

60000/60000 [=====] - 79s 1ms/step - loss: 0.0189 - accuracy: 0.9938 - val\_loss: 0.0210 - val\_accuracy: 0.9929

```

Epoch 9/12
60000/60000 [=====] - 80s 1ms/step - loss: 0.0166 - accuracy: 0.9948 - val_loss: 0.0256 - val_accuracy: 0.9926
Epoch 10/12
60000/60000 [=====] - 80s 1ms/step - loss: 0.0152 - accuracy: 0.9949 - val_loss: 0.0268 - val_accuracy: 0.9918
Epoch 11/12
60000/60000 [=====] - 80s 1ms/step - loss: 0.0136 - accuracy: 0.9954 - val_loss: 0.0235 - val_accuracy: 0.9924
Epoch 12/12
60000/60000 [=====] - 80s 1ms/step - loss: 0.0124 - accuracy: 0.9960 - val_loss: 0.0179 - val_accuracy: 0.9943
Test loss: 0.017945763700317457
Test accuracy: 0.9943000078201294

```

```

In [20]: score = model.evaluate(x_train, y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n*****\n')
#test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
# plot
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch');
ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,12+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Train score: 0.0033918453480264966
Train accuracy: 99.91166591644287

```

```


*****

```

```

Test score: 0.017945763700317457
Test accuracy: 99.43000078201294

```



## CONCLUSION:

```
In [21]: from prettytable import PrettyTable
tb = PrettyTable()
tb.field_names= ("conv_layers", "MAxPoll_layers", "Dense_layers","Dropout", "Accuracy")
tb.add_row(["2", "2", "3", "NO", 98.71])
tb.add_row(["3", "3", "2", "NO", 98.51])
tb.add_row(["4", "2", "2", "NO", 99.29])
tb.add_row(["2", "2", "3", "0.5", 99.12])
tb.add_row(["3", "3", "2", "0.9", 97.47])
tb.add_row(["4", "2", "2", "0.3", 99.43])

print(tb.get_string(titles = "CNN Models - Observations"))
```

conv_layers	MAxPoll_layers	Dense_layers	Dropout	Accuracy
2	2	3	NO	98.71
3	3	2	NO	98.51
4	2	2	NO	99.29
2	2	3	0.5	99.12
3	3	2	0.9	97.47
4	2	2	0.3	99.43

**All the 3 different architectures performed good with accuracy of 98 % Plus and it is also observed that regularizers like drop out resulted in**

**accuracy of 99 % plus**